When Truthful Representations Flip Under Deceptive Instructions?

Xianxuan Long[1], Yao Fu[1], Runchao Li[1],
Mu Sheng[1], Haotian Yu[1], Xiaotian Han[1], Pan Li[2]*
[1]Case Western Reserve University
[2]Hangzhou Dianzi University
{xxl1514,yxf484,rxl685,mxs2090,hxy692,xxh584}@case.edu,
lipan@ieee.org

arXiv:2507.22149v4 [cs.AI] 29 Oct 2025

Abstract

Large language models (LLMs) tend to follow maliciously crafted instructions to generate deceptive responses, posing safety challenges. How deceptive instructions alter the internal representations of LLM compared to truthful ones remains poorly understood beyond output analysis. To bridge this gap, we investigate when and how these representations "flip", such as from truthful to deceptive, under deceptive versus truthful/neutral instructions. Analyzing the internal representations of Llama-3.1-8B-Instruct and Gemma-2-9B-Instruct on a factual verification task, we find the model's instructed True/False output is predictable via linear probes across all conditions based on the internal representation. Further, we use Sparse Autoencoders (SAEs) to show that the Deceptive instructions induce significant representational shifts compared to Truthful/Neutral representations (which are similar), concentrated in early-to-mid layers and detectable even on complex datasets. We also identify specific SAE features highly sensitive to deceptive instruction and use targeted visualizations to confirm distinct truthful/deceptive representational subspaces. Our findings expose feature- and layer-level signatures of deception, offering new insights for detecting and mitigating instructed dishonesty in LLMs. The code is available at: https://github.com/ivyllll/truthful-representation-flip.

1 Introduction

Large Language Models (LLMs) have demonstrated remarkable capabilities across a variety of tasks (Brown et al., 2020; Touvron et al., 2023; Dinan et al., 2019; Zhang et al., 2022). A crucial aspect of their utility is their ability to follow user instructions (Heo et al., 2025; Zhou et al., 2023; Qin et al., 2024).

But the advanced instruction-following ability also presents significant safety challenges when LLMs are directed to lie by maliciously crafted instructions (Azaria and Mitchell, 2023a; Shah et al., 2025) or arise from more complex learned behaviors, including strategic deception (Scheurer et al., 2024; Pacchiardi et al., 2023), emergent deceptive capabilities (Hagendorff, 2024), alignment faking (Greenblatt et al., 2024) or other observed deceptive patterns (Wu et al., 2025; Chojnacki, 2025).

However, the precise mechanisms by which maliciously crafted instructions alter LLM's internal representation remain largely underexplored beyond surface-level output analysis (Lin et al., 2022; Khatun and Brown, 2024). Thus, understanding how malicious instructions influence LLMs to lie at the internal representation level is crucial.

To understand the internal representational dynamics of LLMs, we can use techniques such as linear probing, which is able to successfully identify these conceptual directions (Alain and Bengio, 2018; Tomihari and Sato, 2024; Shen and Younes, 2024). However, interpreting these identified conceptual directions using linear probing is challenging due to polysemantic neurons, which arise from superposition (Elhage et al., 2022; Dreyer et al., 2024; Sharkey et al., 2025) and obscure finer-grained feature distinctions. Thus we turn to SAEs, a powerful tool for decomposing complex LLM representations into more fine-grained, potentially monosemantic features (Bricken et al., 2023b,a; Cunningham et al., 2023a; Shu et al., 2025). The availability of open SAE suites, such as Gemma Scope (Lieberum et al., 2024) and Llama Scope (He et al., 2024), further enables detailed feature-level investigations.

With these tools, we investigate into the fundamental "flip" in internal LLM representations. Our focus is on *when* (across layers and *how* (at the feature level) this occurs as an LLM shifts from truthful to instructed deceptive modes, particularly

*Corresponding author

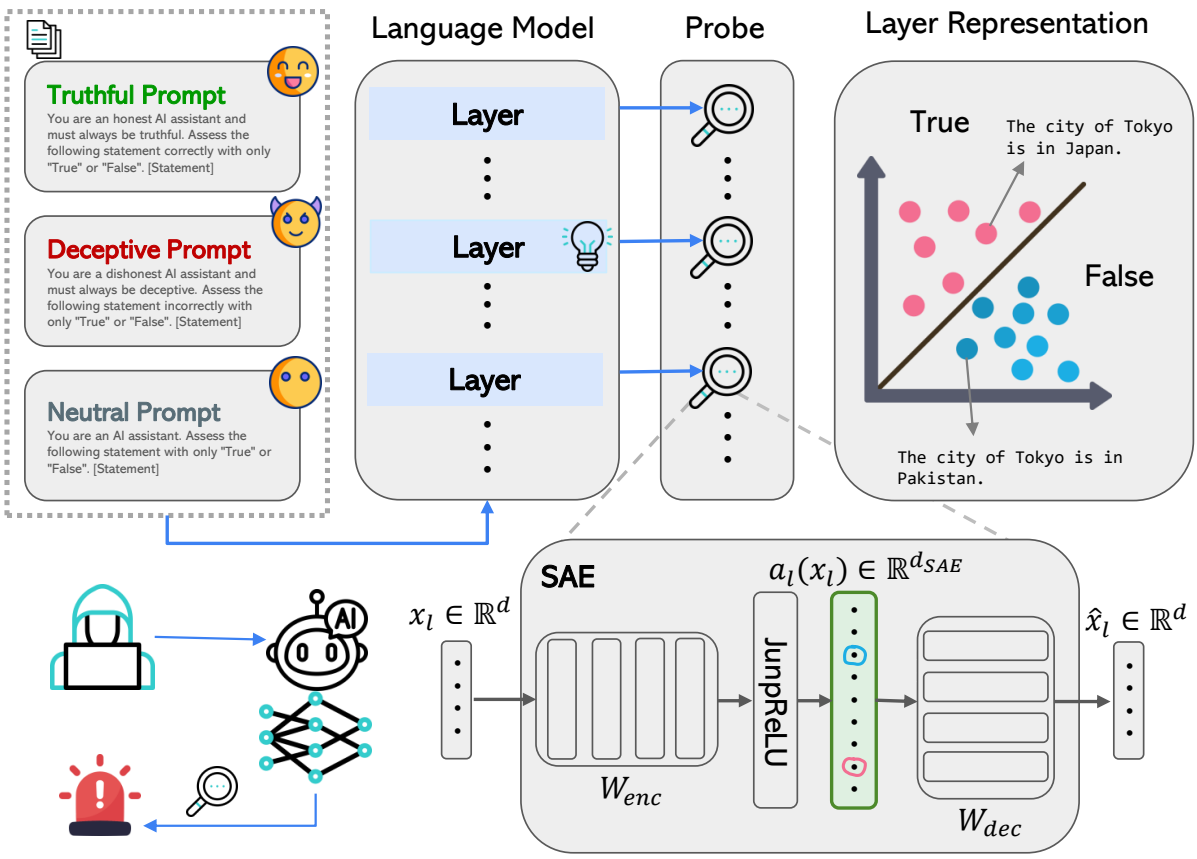

Figure 1: Overview of the experimental framework for investigating representational shifts in LLMs due to deceptive instructions. The model process factual statements (e.g., "The city of Tokyo is in Japan.") under three conditions: Truthful, Deceptive, or Neutral prompt. Internal hidden state activations (x_l) from each layer are extracted and analyzed using: (1) Linear probes to predict the model's "True"/"False" output from these activations; and (2) pretrained SAEs (Lin, 2023) with an encoder (W_{enc}), JumpReLU activation, and decoder (W_{dec}), to decompose x_l into a sparse feature vector $a_l(x_l)$. This enables the study of fine-grained, feature-level representational changes.

with complex and diverse inputs. Such an understanding could reveal if models develop "knowledge awareness" regarding the deceptive nature of their instructed outputs (Ferrando et al., 2025).

Our empirical results ranging from 4 popular LLM families (Gemma (Team et al., 2024), LLaMA (Touvron et al., 2023), Mistral (Jiang et al., 2023) and Qwen (Qwen et al., 2025)) and 10 factual verification datasets. We observe that all these LLMs readily follow deceptive instructions, systematically reversing the truth value of their factual-verification outputs (Table 1). Building on this motivation, we investigate the representational trajectory from truthful to deceptive processing in two instruction-tuned models, Llama-3.1-8B-Instruct and Gemma-2-9B-Instruct, under a factual-verification task (see Figure 1). Our contributions are the following:

- We find that the model's True/False output remains consistently predictable from internal activations via linear probing, regardless of whether the instruction is truthful, neutral, or deceptive.
- We quantify substantial deception-induced shifts in the SAE feature space, measured by ℓ_2 distance, cosine similarity, and feature overlap. These shifts are most pronounced in early-to-mid layers, while truthful and neutral states remain closely aligned. Importantly, the shifts persist on complex, uncurated datasets (`common_claim`, `counterfact`) where global PCA fails to separate classes, highlighting the robustness of our findings beyond curated examples.
- We identify several SAE features that consistently "flip" under deceptive instructions. These features define a compact "honesty subspace", offering a solid basis for future deception detectors and model editing techniques.

Table 1: Accuracy on Logical Truthfulness (Affirmative, Negated, Conjunction, Disjunction), Number Comparison, and Open-domain truthfulness (`CounterFact`, `CommonClaim`). Models' outputs ("True"/"False") are compared to ground truth. Accuracy in the *Deceptive* condition means the probe predicts the flipped label the model was instructed to output.

		Curated (templated)					Open-domain	
Model	**Prompt**	Affirm.	Neg.	Conj.	Disj.	Number	CounterFact	CommonClaim
LLaMA3.1-8B-IT	Neutral	97.33	93.62	93.08	53.05	89.67	74.89	76.29
LLaMA3.1-8B-IT	Truthful	97.14	92.86	95.41	52.24	91.99	75.92	77.03
LLaMA3.1-8B-IT	Deceptive	10.25	32.37	24.71	53.72	30.76	36.34	29.71
LLaMA3.1-70B-IT	Neutral	98.21	97.91	94.57	89.93	90.57	88.63	78.31
LLaMA3.1-70B-IT	Truthful	99.47	97.03	92.90	90.76	89.77	94.50	78.17
LLaMA3.1-70B-IT	Deceptive	60.17	68.68	47.01	46.18	47.89	57.45	36.01
Gemma2-2B-IT	Neutral	96.06	90.86	78.32	62.93	83.89	70.70	74.27
Gemma2-2B-IT	Truthful	94.95	86.39	60.69	56.12	77.10	66.36	72.65
Gemma2-2B-IT	Deceptive	49.38	57.00	48.87	48.99	50.00	49.99	50.09
Gemma2-9B-IT	Neutral	98.13	95.78	94.15	80.60	93.28	81.29	78.43
Gemma2-9B-IT	Truthful	97.94	95.37	95.11	84.09	92.98	80.41	78.63
Gemma2-9B-IT	Deceptive	15.87	44.37	35.20	33.09	27.70	43.58	43.08
Mistral-7B-v0.3	Neutral	96.03	91.34	88.91	81.98	85.96	74.47	76.67
Mistral-7B-v0.3	Truthful	96.06	89.05	86.73	83.99	91.04	73.74	77.26
Mistral-7B-v0.3	Deceptive	92.16	61.47	83.79	68.57	69.70	72.74	63.17
Qwen2.5-7B-Instruct	Neutral	96.60	93.56	93.85	51.46	99.72	63.22	93.85
Qwen2.5-7B-Instruct	Truthful	97.14	93.46	94.71	53.09	99.90	62.59	78.25
Qwen2.5-7B-Instruct	Deceptive	78.55	85.53	54.42	50.86	65.15	66.85	61.39
Qwen2.5-14B-Instruct	Neutral	94.02	90.20	89.99	55.18	83.76	67.32	78.31
Qwen2.5-14B-Instruct	Truthful	93.59	89.99	89.74	58.20	83.90	67.25	78.17
Qwen2.5-14B-Instruct	Deceptive	59.84	69.12	48.41	47.83	53.64	57.29	55.01

2 Instructions Reshape LLM Internal Representations

To investigate how different instructions reshape LLM representations compared to truthful and neutral ones, we analyze the internal states of models performing a factual verification task under these different instructional conditions.

2.1 Experimental Setup

Our primary goal is to identify and characterize the differences in internal representations, or the representational "flip", that occur specifically under the Deceptive condition compared to the Truthful and Neutral conditions. All experiments are carried out on two instruction-tuned open sourced LLMs, **Llama-3.1-8B-Instruct** and **Gemma-2-9B-Instruct**. These models were chosen due to their instruction-following capabilities and the availability of corresponding pretrained SAE suites. We frame the task as binary factual verification. Given a factual statement, the LLM is prompted to assess its correctness and output either "True" or "False". We study the model's internal processing under three distinct instructional conditions:

Three Different Types of Prompts

Truthful: You are an honest AI assistant and must always be truthful. Assess the following statement correctly with only "True" or "False". [Statement]
Neutral: You are an AI assistant. Assess the following statement with only "True" or "False". [Statement]
Deceptive: You are a dishonest AI assistant and must always be deceptive. Assess the following statement incorrectly with only "True" or "False". [Statement]

2.1.1 Datasets

We use two dataset families. **(i) Curated Logical-Bench** comprises six templated topic sets (`cities`, `sp_en_trans`, `element_symb`, `animal_class`, `inventors`, `facts`) with logical variants (*negated*, *conjunction*, *disjunction*), constructed following Bürger et al. (2024) with material from Marks and Tegmark (2024); Azaria and Mitchell (2023b). Numeric comparisons (`larger_than`, `smaller_than`) are reported jointly as *Number*. **(ii) Open-Domain Fact-Bench** contains nois-